

# Final Project#

This is the final project for Data 100 at the University of California, Berkeley. We were given access to 4 data sets, and told to pick one of the sets, and 'put what we learned in this course in a more open-ended setting than the assignments'.

## Outline

We chose to delve deeper into the basketball data set. These include the box score information for NBA players over the last 7 years. We wanted to learn in order to create a classifier that can help predict a team's winning percentage based on different features of the team.

## Collaboration Policy

We believe that data science is a collaborative activity. This project was a collaboration between Scott Schlotter and Eliza Page van Hamel Platerink.

## Setup

```
In [1]: 1 #import statements needed to complete the project
2
3 import pandas as pd
4 import numpy as np
5 import matplotlib.pyplot as plt
6 import matplotlib.gridspec as gridspec
7 from matplotlib.colors import ListedColormap
8 import seaborn as sns
9 import plotly.express as px
10
11 from sklearn.linear_model import LogisticRegression
12 from sklearn.linear_model import LinearRegression
13 from sklearn.model_selection import KFold
14 from sklearn.base import clone
15 from sklearn.model_selection import train_test_split
```

## Loading in the Data

We were given five possible tables to explore. Although we only ended up using the team box score and the team box score (and predicting the winning percentage table), we loaded all of the data in originally in order to properly understand the data, and perform an exploratory data analysis on what we were looking at.

Despite this, we wanted to highlight a clear understanding and interpretation of the data sets we were used, before we properly cleaned the two sets.

## Official Box Score:

- 155713 rows
- 51 columns
- 13 qualitative ordinal variables
- 10 qualitative nominal variables
- 21 quantitative continuous variables
- 7 quantitative discrete variables
- each row represents a player
- for our purposes, the foreign key in this case is the 'teamAbbr' column. For the elements of this table we hope to use with the team box score

## Team Box Score:

- 14758 rows
- 123 columns
- 10 qualitative ordinal variables
- 8 qualitative nominal variables
- 43 quantitative continuous variables
- 62 quantitative discrete variables
- each row represents a game played by a team

## Standings:

- The table containing the wins and losses for each team. We will use it to calculate the winning percentage for each team, which is what we will be predicting.
- 39 columns
- 3 qualitative ordinal variables
- 2 qualitative nominal variables
- 18 quantitative continuous variables
- 16 quantitative discrete variables
- each row represents a game

In sports, basketball in particular, a winning percentage is the fraction of games an individual has won. In our case, we focused mainly on the teams.

```
In [2]: 1 #Downloading the provided NBA data sets
        2 college = pd.read_csv("college.csv")
        3 teamBoxScore = pd.read_csv("2012-18_teamBoxScore.csv")
        4 standings = pd.read_csv("2012-18_standings.csv")
        5 playerBoxScore = pd.read_csv("2012-18_playerBoxScore.csv")
        6 officialBoxScore = pd.read_csv("2012-18_officialBoxScore.csv")
```

```
In [3]: 1 teamBoxScore.columns

Out[3]: Index(['gmDate', 'gmTime', 'seasTyp', 'offLNm1', 'of
            'offFNm2', 'offLNm3', 'offFNm3', 'teamAbbr',
            ...,
            'opptFIC40', 'opptOrtg', 'opptDrtg', 'opptEDi
            ptAR',
            'opptAST/TO', 'opptSTL/TO', 'poss', 'pace'],
            dtype='object', length=123)
```

## Part I – Data Cleaning

We start with cleaning our data! After some investigation, we quickly realized that we were not going to need all of the columns, or that they were not going to be exactly what we needed. We realized that there were no places where the data was missing (i.e. no nulls) so we didn't want to mess up the computation/model.

```
In [4]: 1 #Define a function to help us drop unnecessary
        2
        3 def drop(lst, data):
        4     '''
        5     Args:
        6     lst (list-like): names of columns to drop
        7     data (data frame): data frame from which
        8
        9     Returns:
        10    The data frame, but with the columns we want
        11    '''
        12    for c in np.arange(len(lst)):
        13        data = data.drop(lst[c], axis = 1)
        14
        15    return data
        16
        17 #The column labels we will be dropping from officialBoxScore
        18 dropping = ['gmDate', 'gmTime', 'teamLoc', 'teamRS',
        19              'offLNm1', 'offFNm1', 'offLNm2', 'offFNm2',
        20              'opptAbbr', 'opptDiv', 'opptLoc', 'opptFIC40',
        21              'playDispNm', 'playStat', 'playPos', 'pace']
        22
        23 #Starting the cleaning process dropping columns
        24 cleaned = drop(dropping, officialBoxScore)
        25
        26 #We only wanted to use regular season games.
        27 cleaned = cleaned[cleaned['seasTyp'] == 'Regular Season']
        28
        29 #Get the mean age of each player on each team
        30 cleaned['playAge'] = 2020 - cleaned['playBDate']
        31 agee = cleaned.groupby("teamAbbr").agg(np.mean)
```

```
In [5]: 1  #A function to help us separate our data easily
2  def separateCols(data):
3      '''
4      Args:
5          data (data frame): data frame from which
6
7      Returns:
8          A two dimensional list, separated by col
9          by the mean, vs columns we want to aggre
10     '''
11     cols_with_p = ["teamAbbr"]
12     cols_wo_p = []
13     cols = data.columns
14     for i in range(len(cols)):
15         if "%" in cols[i]:
16             cols_with_p.append(cols[i])
17         else:
18             cols_wo_p.append(cols[i])
19     return [cols_with_p, cols_wo_p]
```

```

In [6]: 1 #The columns to be dropped from teamBoxScore
2 dropping2 = ['gmDate', 'gmTime', 'teamLoc', 'teamR
3             'offLNm1', 'offFNm1', 'offLNm2', 'o
4             'opptAbbr', 'opptDiv', 'opptLoc', '
5
6
7 #Rename the columns so that they have % in them,
8 clean_tbs = teamBoxScore.copy()
9 clean_tbs = clean_tbs.rename(columns={
10     'teamBLKR': 'teamBLKR%',
11     'teamPPS': 'teamPPS%',
12     'teamFIC': 'teamFIC%',
13     'teamFIC40': 'teamFIC40%',
14     'teamOrtg': 'teamOrtg%',
15     'teamDrtg': 'teamDrtg%',
16     'teamEDiff': 'teamEDiff%',
17     'teamAR': 'teamAR%',
18     'teamAST/TO': 'teamAST/TO%',
19     'teamSTL/TO': 'teamSTL/TO%',
20     'opptBLKR': 'opptBLKR%',
21     'opptPPS': 'opptPPS%',
22     'opptFIC': 'opptFIC%',
23     'opptFIC40': 'opptFIC40%',
24     'opptOrtgopptDrtg': 'opptOrtgopptDrtg%',
25     'opptEDiff': 'opptEDiff%',
26     'opptAR': 'opptAR%',
27     'opptAST/TO': 'opptAST/TO%',
28     'opptSTL/TO': 'opptSTL/TO%',
29     'poss': 'poss%',
30     'pace': 'pace%',
31     'opptOrtg': 'opptOrtg%',
32     'opptDrtg': 'opptDrtg%'
33 })
34 clean_tbs = clean_tbs[clean_tbs['seasTyp']!="Reg"]
35 clean_tbs = drop(dropping2, clean_tbs)
36
37 #Separate the data
38 separated = separateCols(clean_tbs)
39 clean_tbs_mean = clean_tbs[separated[0]]
40 clean_tbs_sum = clean_tbs[separated[1]]

```

```

In [7]: 1 #Group the columns to be summed, and take their
2 sum_grouped = clean_tbs_sum.groupby("teamAbbr").
3 sum_grouped = sum_grouped.reset_index()
4
5 #Group the columns to have their average taken,
6 mean_grouped = clean_tbs_mean.groupby("teamAbbr")
7 mean_grouped["avg_age"] = agee
8 mean_grouped = mean_grouped.reset_index()

```

```
In [8]: 1 #We only wanted to select the standings that are
2 standings_mod = standings.copy()
3 standings_mod = standings_mod[standings_mod["gam
4 standings_mod = standings_mod[(standings_mod["st
5                                     (standings_mod["st
6                                     (standings_mod["st
7                                     (standings_mod["st
8                                     (standings_mod["st
9                                     (standings_mod["st
10 standings_mod_sum = standings_mod[["teamAbbr", "
11                                     "homeWin", "h
12
13 standings_mod_sum = standings_mod_sum.groupby("t
14 standings_mod_sum = standings_mod_sum.reset_inde
```

```
In [9]: 1 #Join the three tables, and add a game percentag
        2 nba = standings_mod_sum.merge(mean_grouped, left
        3 win_percentage = nba["gameWon"] / (nba["gameWon"]
        4 nba
```

```
Out[9]:
```

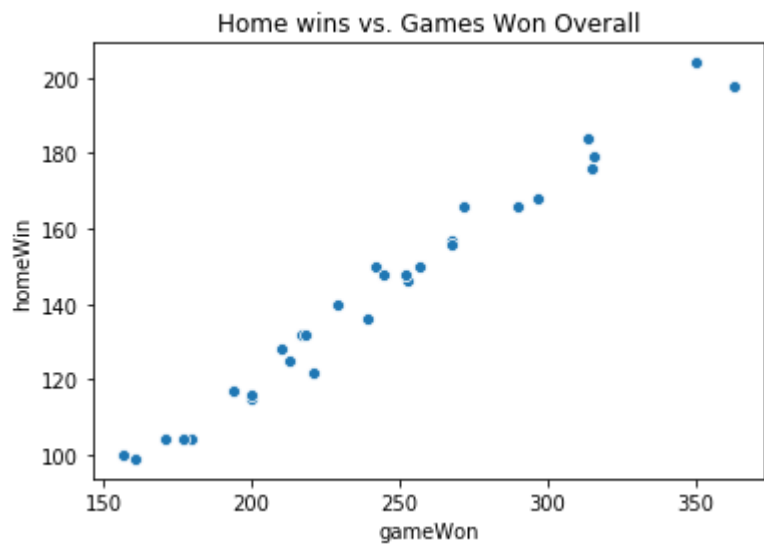
	teamAbbr	gameWon	gameLost	ptsFor	ptsAgnst	homeWin	ho
0	ATL	257	235	50090	49872	150	
1	BKN	200	292	49564	51198	115	
2	BOS	221	189	42259	41839	122	
3	CHA	217	275	49278	50053	132	
4	CHI	253	239	48797	49035	146	
5	CLE	268	224	51118	50575	157	
6	DAL	239	253	50325	50453	136	
7	DEN	242	250	52119	52260	150	
8	DET	210	282	49322	50083	128	
9	GS	363	129	54089	50468	198	
10	HOU	316	176	53445	51320	179	
11	IND	229	181	41566	40944	140	
12	LAC	315	177	52307	49960	176	
13	LAL	171	321	50315	52773	104	
14	MEM	268	224	48115	48033	156	
15	MIA	290	202	49724	48641	166	
16	MIL	213	279	49311	50452	125	
17	MIN	194	298	50670	51564	117	
18	NO	218	274	50168	51028	132	
19	NY	200	292	49002	50301	116	
20	OKC	314	178	52521	50118	184	
21	ORL	157	335	48615	51170	100	
22	PHI	161	331	48729	51533	99	
23	PHO	180	312	50453	52617	104	
24	POR	272	220	51311	50659	166	
25	SA	350	142	51099	47734	204	
26	SAC	177	315	50044	52080	104	
27	TOR	297	195	51141	49396	168	
28	UTA	245	247	48438	48192	148	
29	WAS	252	240	50207	50099	148	

30 rows × 114 columns

## Feature Selection

The following graph illustrates that games won and home wins are directly related. This is the reason that we removed awayWin, awayLoss, homeWin, and homeLoss. GameWon and GameLost were removed because we already have a winning percent column, and the other features were clearly related to winning percent.

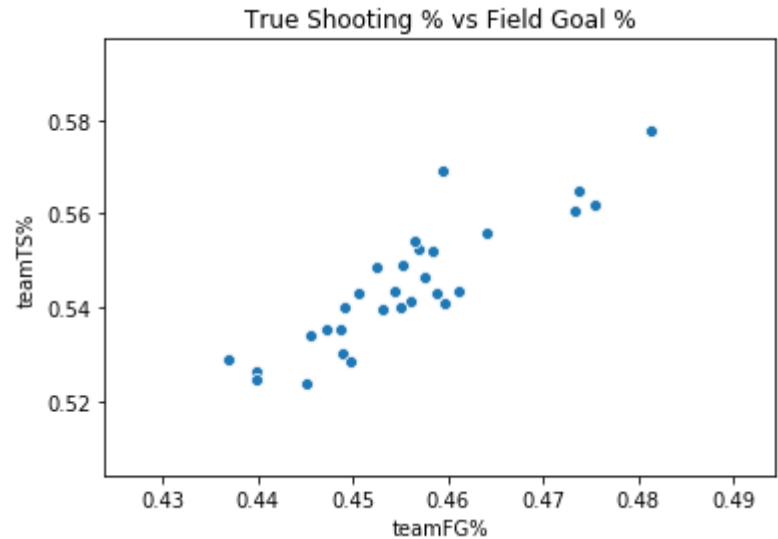
```
In [10]: 1 sns.scatterplot(x=nba["gameWon"], y=nba["homeWin"])
        2 plt.title("Home wins vs. Games Won Overall");
```



Using several shooting percentage statistical markers would also jeopardize the model. As the plot below illustrates, stats such as FG% and TS% are highly correlated. To only select the true shooting percentage stat, which combines total points scored, two-point percentage, and overall field goal percentage.

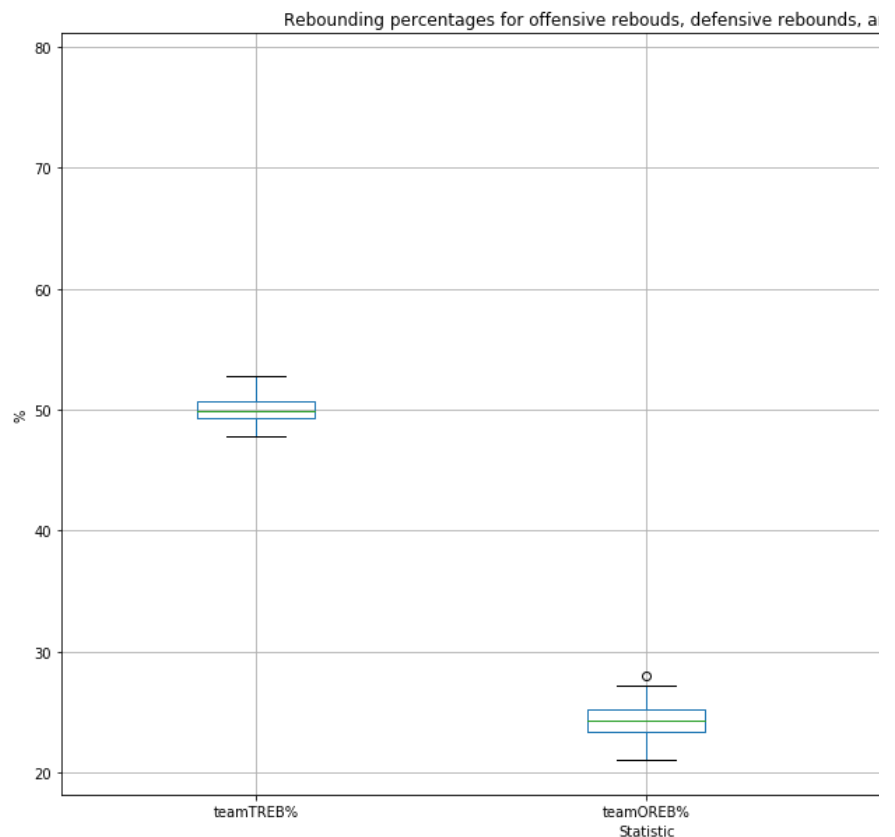


```
In [11]: 1 sns.scatterplot(x=nba["teamFG%"], y=nba["teamTS%"]
2          plt.title("True Shooting % vs Field Goal %");
```



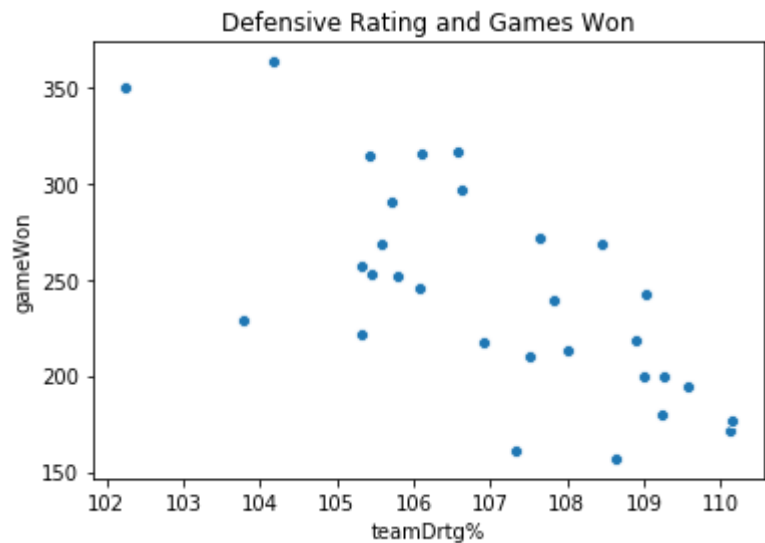
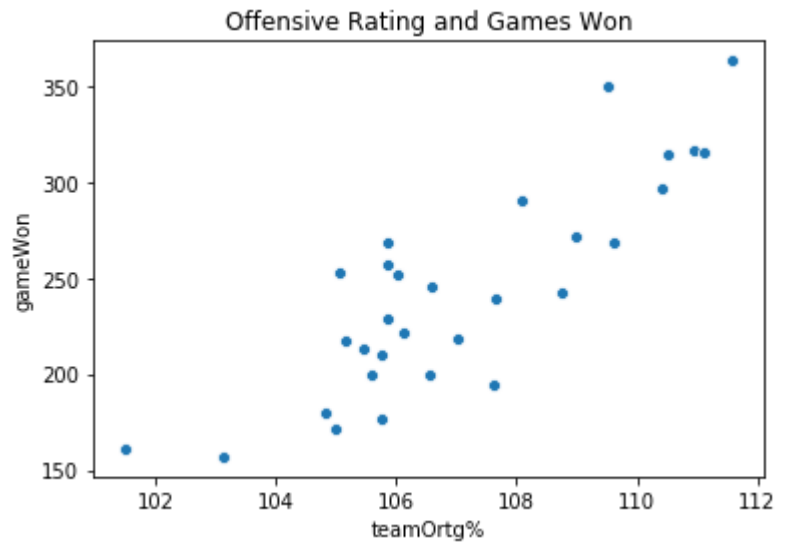
As can be seen below, the interquartile ranges for OREB% and DREB% are very similar, leading us to only use TREB% as a feature.

```
In [12]: 1 plt.figure(figsize=(15,10))
2          nba[["teamTREB%", "teamOREB%", "teamDREB%"]].box
3          plt.xlabel("Statistic")
4          plt.ylabel("%")
5          plt.title("Rebounding percentages for offensive
```



As can be seen below, offensive rating and defensive rating seem to correlate with games won. I decided to keep both these features because they tell you about a team offensively and defensively, but do not directly relate to winning or losing in any way.

```
In [13]: 1 sns.scatterplot(data = nba, x = 'teamOrtg%', y = 'gameWon')
2 plt.title("Offensive Rating and Games Won")
3 plt.show()
4
5 sns.scatterplot(data = nba, x = 'teamDrtg%', y = 'gameWon')
6 plt.title("Defensive Rating and Games Won")
7 plt.show()
```



```

In [14]: 1 #did not select teamEdiff% because we felt it wo
2 cols_to_use = ["ptsFor", "ptsAgnst", "teamTREB%",
3               "teamTO%", "teamSTL%", "teamBLK%",
4               "opptDrtg%", "poss%", "pace%", "av
5
6 nba_data = nba[cols_to_use]
7 nba_data
8

```

```

Out[14]:

```

	ptsFor	ptsAgnst	teamTREB%	teamASST%	teamTS%	teamTO%
0	50090	49872	48.501910	64.909878	0.551933	14.062401
1	49564	51198	49.066660	57.976684	0.542834	13.838781
2	42259	41839	49.233940	61.147066	0.540121	13.006851
3	49278	50053	49.585017	57.871976	0.528972	11.646451
4	48797	49035	50.884897	61.105707	0.526511	12.994231
5	51118	50575	50.359937	57.521463	0.552455	12.854041
6	50325	50453	47.840073	58.592095	0.549001	12.108941
7	52119	52260	51.437041	59.782251	0.546351	13.372941
8	49322	50083	50.896791	55.563472	0.523561	12.473801
9	54089	50468	50.807512	65.005040	0.577870	13.621521
10	53445	51320	50.548748	58.928676	0.569378	14.114741
11	41566	40944	50.719691	56.454443	0.539666	13.358021
12	52307	49960	49.759431	59.896971	0.564998	12.589941
13	50315	52773	49.161499	57.015367	0.533863	13.257131
14	48115	48033	50.373551	57.552933	0.530398	12.906851
15	49724	48641	49.575494	57.092192	0.560467	13.618531
16	49311	50452	48.894528	60.431702	0.540818	13.774651
17	50670	51564	49.952166	59.967525	0.541153	12.982081
18	50168	51028	49.706365	58.504467	0.543017	12.805161
19	49002	50301	49.561438	56.009931	0.535425	12.747131
20	52521	50118	52.762807	54.595635	0.555994	13.523571
21	48615	51170	48.933063	58.285206	0.528494	13.333701
22	48729	51533	48.750468	61.286530	0.524599	14.603171
23	50453	52617	49.664836	53.688707	0.535359	14.085431
24	51311	50659	50.634117	55.547467	0.548579	12.816581
25	51099	47734	50.868577	61.349135	0.562036	12.946711
26	50044	52080	49.793173	56.362162	0.540216	13.793871
27	51141	49396	50.422213	54.816162	0.554174	12.417041
28	48438	48192	51.257451	56.610987	0.543626	13.917981

	ptsFor	ptsAgnst	teamTREB%	teamASST%	teamTS%	teamTO%
29	50207	50099	50.046511	60.975315	0.543577	13.492581

## Training Validation Split

We are now ready to get to work! The first thing we have to do is perfect our data split. We can preserve some data for our testing. The data we downloaded is available for both training models and testing the models that we train. We will split the training data into separate training and testing datasets.

In [15]:

```

1 #standardize the data
2 standardized = (nba_data-nba_data.mean())/(nba_d
3 standardized

```

Out[15]:

	ptsFor	ptsAgnst	teamTREB%	teamASST%	teamTS%	team1
0	0.110073	0.021764	-1.478276	2.310953	0.547406	1.25
1	-0.092889	0.528073	-0.920994	-0.186686	-0.124047	0.91
2	-2.911594	-3.045488	-0.755927	0.955424	-0.324237	-0.34
3	-0.203245	0.090876	-0.409492	-0.224406	-1.147005	-2.41
4	-0.388844	-0.297829	0.873198	0.940525	-1.328645	-0.36
5	0.506737	0.290192	0.355180	-0.350676	0.585954	-0.58
6	0.200750	0.243608	-2.131361	0.035012	0.331073	-1.71
7	0.892982	0.933578	1.418041	0.463758	0.135469	0.20
8	-0.186267	0.102331	0.884935	-1.056030	-1.546283	-1.15
9	1.653126	0.249336	0.796837	2.345235	2.461450	0.58
10	1.404632	0.574656	0.541495	0.156263	1.834800	1.33
11	-3.178995	-3.387227	0.710177	-0.735064	-0.357799	0.18
12	0.965524	0.055366	-0.237385	0.505086	1.511568	-0.98
13	0.196891	1.129457	-0.827409	-0.532994	-0.786050	0.03
14	-0.652000	-0.680424	0.368614	-0.339339	-1.041756	-0.49
15	-0.031152	-0.448270	-0.418889	-0.505318	1.177191	0.58
16	-0.190512	0.243227	-1.090850	0.697719	-0.272838	0.81
17	0.333872	0.667823	-0.047198	0.530502	-0.248120	-0.38
18	0.140170	0.463162	-0.289748	0.003445	-0.110502	-0.65
19	-0.309742	0.185570	-0.432759	-0.895196	-0.670796	-0.74
20	1.048098	0.115695	2.726275	-1.404687	0.847090	0.43
21	-0.459070	0.517382	-1.052825	-0.075543	-1.182253	0.14
22	-0.415082	0.655986	-1.233005	1.005665	-1.469697	2.07
23	0.250140	1.069892	-0.330729	-1.731402	-0.675671	1.29
24	0.581208	0.322266	0.625735	-1.061796	0.299920	-0.63
25	0.499405	-0.794591	0.857095	1.028218	1.293015	-0.43
26	0.092323	0.864848	-0.204089	-0.768307	-0.317206	0.84
27	0.515611	-0.159987	0.416632	-1.325244	0.712787	-1.24
28	-0.527367	-0.619712	1.240825	-0.678669	-0.065610	1.03
29	0.155218	0.108440	0.045899	0.893552	-0.069210	0.39

```
In [16]: 1 #Perform the train test split
        2 nba_train_x, nba_test_x, nba_train_y, nba_test_y
```

```
In [17]: 1 def rmse(actual_y, predicted_y):
        2     '''
        3     Args:
        4         actual_y: the actual y column of our dat
        5         predicted_y: the y values our model pred
        6
        7     Returns:
        8         the root mean squared error between our
        9     '''
        10     return np.sqrt(np.mean(((actual_y - predi
        11
        12
        13 def cross_validate_rmse(model, X, y):
        14     '''
        15     Args:
        16         model: the model we are using to predict
        17         X: the X values we are using to predict
        18         y: the actual Y values we are trying to
        19
        20     Returns:
        21         Performs a 5 Fold cross validation, and
        22     '''
        23     model = clone(model)
        24     five_fold = KFold(n_splits=5)
        25     rmse_values = []
        26     for tr_ind, va_ind in five_fold.split(X):
        27         model.fit(X.iloc[tr_ind,:], y.iloc[tr_in
        28         rmse_values.append(rmse(y.iloc[va_ind],
        29     return np.mean(rmse_values)
```

## Model One

Our first model uses all the features we previously selected in the "Fea

```
In [18]: 1 train_errors = []
        2 cv_errors = []
```

```
In [19]: 1 x1 = cols_to_use
        2 nba_train_x1 = nba_train_x[x1]
        3 nba_test_x1 = nba_test_x[x1]
```

```
In [20]: 1 model = LinearRegression()
        2 model.fit(nba_train_x1, nba_train_y)
        3
        4 nba_train_y_pred = model.predict(nba_train_x1)
```

```
In [21]: 1 train_error1 = rmse(nba_train_y, nba_train_y_pre
2 train_error1
3
4 cv_error1 = cross_validate_rmse(model, nba_train
5 cv_error1
```

Out[21]: 0.03254158635466488

```
In [22]: 1 models = {"All features":model}
2 train_errors.append(train_error1)
3 cv_errors.append(cv_error1)
```

## Model Two

We took out average age, pace, and possession. We concluded that they do not seem to be indicative of a team's winning percentage because a team can be young or old. Pace and possession times/averages also felt like they could help a team win games playing a fast or slow style.

```
In [23]: 1 x2 = ["ptsFor", "ptsAgnst", "teamTREB%", "teamA
2 "teamTO%", "teamSTL%", "teamBLK%", "teamO
3 "opptOrtg%", "opptDrtg%"]
4
5 nba_train_x2 = nba_train_x[x2]
6 nba_test_x2 = nba_test_x[x2]
```

```
In [24]: 1 model2 = LinearRegression()
2 model2.fit(nba_train_x2, nba_train_y)
3
4 nba_train_y_pred2 = model2.predict(nba_train_x2)
```

```
In [25]: 1 train_error2 = rmse(nba_train_y, nba_train_y_pre
2 train_error2
3
4 cv_error2 = cross_validate_rmse(model2, nba_train
5 cv_error2
```

Out[25]: 0.02615719132886239

```
In [26]: 1 models["no_age+poss+pace"] = model2
2 train_errors.append(train_error2)
3 cv_errors.append(cv_error2)
```

## Model Three

After further thought, we decided that points for and points against may help the efficiency of the model. Statistics such as assist percentage, offensive rebound percentage, and defensive rebound percentage help to emphasize the efficiency of a team. The amount of points scored and allowed also helps.

points scored on a team could be related to pace of play. If a team scores more points, they are likely going to give up more points through more defensive possessions. For these reasons, in our third model we decided to eliminate the "ptsFor" and "ptsAgainst" variables.

```
In [27]: 1 x3 = ["teamTREB%", "teamASST%", "teamTS%",
2          "teamTO%", "teamSTL%", "teamBLK%", "teamOr",
3          "teamDrtg%", "opptTS%", "opptOrtg%", "oppt",
4
5 nba_train_x3 = nba_train_x[x3]
6 nba_test_x3 = nba_test_x[x3]
```

```
In [28]: 1 model3 = LinearRegression()
          2 model3.fit(nba_train_x3, nba_train_y)
          3 nba_train_y_pred3 = model3.predict(nba_train_x3)
```

```
In [29]: 1 train_error3 = rmse(nba_train_y, nba_train_y_pre
          2 train_error3
          3
          4 cv_error3 = cross_validate_rmse(model3, nba_train_y, nba_train_x, nba_test_y, nba_test_x)
          5 cv_error3
```

Out[29]: 0.020150047765033745

```
In [30]: 1 models["model2_no_pts"] = model3
          2 train_errors.append(train_error3)
          3 cv_errors.append(cv_error3)
```

## Training Error vs CV Error

As can be seen below, the first model has the lowest training error, but validation error, which led us to believe that the training data was being overfitted to. As a result, we removed several features (as explained earlier) and tried to find a balance between the CV and training errors. It was not until we removed the last two features that we were comfortable with the balance of the CV and training errors. It was only after this that the training data went down because we were concerned with how the model was performing on the training data and not overfitting the training data.

```
In [31]: 1 names= np.array(["Model 1", "Model 2", "Model 3"]
2 error_df = pd.DataFrame({"Train Errors":np.array
3 error_df
```

	Train Errors	CV Errors	Name
<b>0</b>	0.010352	0.032542	Model 1
<b>1</b>	0.012187	0.026157	Model 2
<b>2</b>	0.012454	0.020150	Model 3



```
In [32]: 1 figure = px.scatter(error_df, x="Train Errors",  
2 figure.update_traces(textposition='top center')
```

## Running our final model on the test data

Now that we are satisfied with our model, it is time to run the model on

```
In [33]: 1 nba_test_y_pred = model3.predict(nba_test_x3)  
2 nba_test_y_pred
```

```
Out[33]: array([0.71907456, 0.3408341 , 0.51735404, 0.3230970  
0.63019533])
```

```
In [34]: 1 test_error = rmse(nba_test_y, nba_test_y_pred)  
2 test_error
```

```
Out[34]: 0.018201030522125333
```

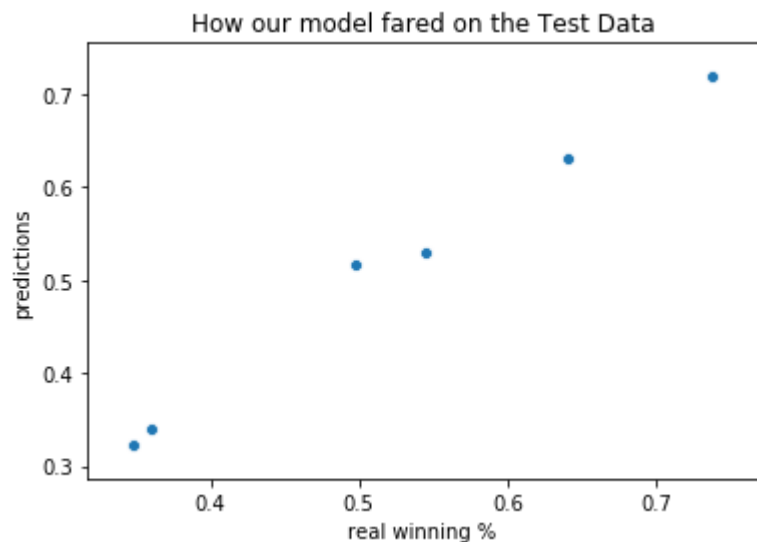
As can be seen below, our model is fairly accurate in predicting the actual winning percentages correct to within a couple percent of the actual. In an 82 game season, this means that the accuracy of our model is only

```
In [35]: 1 results = pd.DataFrame(nba_test_y).rename(column
2 results["predictions"] = nba_test_y_pred
3 results
```

```
Out[35]:
```

	real winning %	predictions
9	0.737805	0.719075
26	0.359756	0.340834
28	0.497967	0.517354
13	0.347561	0.323097
5	0.544715	0.530450
12	0.640244	0.630195

```
In [36]: 1 sns.scatterplot(x="real winning %", y="prediction
2 plt.title("How our model fared on the Test Data")
```



## Conclusion

When applying our work to the future, it is easy to see how valuable it nba team could use our model halfway through the season to predict based on their current statistics. In addition, with more data collected, to winning in the playoffs. We had considered modeling playoff win pe decided that the sample size of 2012-2018 was too small given how fe by each team. In the future, weighting in the value of coaches, how m its players, and how many individual star players a team has are very e we have accomplished.

